

# km5188\_ds4e\_hw5

December 8, 2022

## 1 DS4E: Homework 4

```
[5]: pip install pandas
```

```
Requirement already satisfied: pandas in /opt/conda/lib/python3.10/site-packages (1.5.2)
Requirement already satisfied: pytz>=2020.1 in /opt/conda/lib/python3.10/site-packages (from pandas) (2022.2.1)
Requirement already satisfied: python-dateutil>=2.8.1 in /opt/conda/lib/python3.10/site-packages (from pandas) (2.8.2)
Requirement already satisfied: numpy>=1.21.0 in /opt/conda/lib/python3.10/site-packages (from pandas) (1.23.2)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.10/site-packages (from python-dateutil>=2.8.1->pandas) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

```
[4]: pip install statsmodels
```

```
Requirement already satisfied: statsmodels in /opt/conda/lib/python3.10/site-packages (0.13.5)
Requirement already satisfied: packaging>=21.3 in /opt/conda/lib/python3.10/site-packages (from statsmodels) (21.3)
Requirement already satisfied: scipy>=1.3 in /opt/conda/lib/python3.10/site-packages (from statsmodels) (1.9.0)
Requirement already satisfied: pandas>=0.25 in /opt/conda/lib/python3.10/site-packages (from statsmodels) (1.5.2)
Requirement already satisfied: patsy>=0.5.2 in /opt/conda/lib/python3.10/site-packages (from statsmodels) (0.5.3)
Requirement already satisfied: numpy>=1.17 in /opt/conda/lib/python3.10/site-packages (from statsmodels) (1.23.2)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /opt/conda/lib/python3.10/site-packages (from packaging>=21.3->statsmodels) (3.0.9)
Requirement already satisfied: python-dateutil>=2.8.1 in /opt/conda/lib/python3.10/site-packages (from pandas>=0.25->statsmodels) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /opt/conda/lib/python3.10/site-packages (from pandas>=0.25->statsmodels) (2022.2.1)
Requirement already satisfied: six in /opt/conda/lib/python3.10/site-packages
```

(from patsy>=0.5.2->statsmodels) (1.16.0)

Note: you may need to restart the kernel to use updated packages.

```
[3]: pip install sklearn
```

Requirement already satisfied: sklearn in /opt/conda/lib/python3.10/site-packages (0.0.post1)

Note: you may need to restart the kernel to use updated packages.

```
[2]: pip install scikit-learn
```

Requirement already satisfied: scikit-learn in /opt/conda/lib/python3.10/site-packages (1.2.0)

Requirement already satisfied: joblib>=1.1.1 in /opt/conda/lib/python3.10/site-packages (from scikit-learn) (1.2.0)

Requirement already satisfied: numpy>=1.17.3 in /opt/conda/lib/python3.10/site-packages (from scikit-learn) (1.23.2)

Requirement already satisfied: scipy>=1.3.2 in /opt/conda/lib/python3.10/site-packages (from scikit-learn) (1.9.0)

Requirement already satisfied: threadpoolctl>=2.0.0 in /opt/conda/lib/python3.10/site-packages (from scikit-learn) (3.1.0)

Note: you may need to restart the kernel to use updated packages.

```
[5]: # the usual libraries
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

# statsmodels
import statsmodels.formula.api as smf

# SciKit Learn libraries
from sklearn.model_selection import train_test_split # for splitting data
from sklearn.linear_model import LinearRegression   # linear regression
from sklearn.preprocessing import StandardScaler     # feature scaling
from sklearn import metrics                         # for evaluation metrics
from sklearn.neighbors import KNeighborsClassifier   # knn
```

## 1.1 Question 1

1(a)

```
[7]: bechdel = pd.read_csv('bechdel.csv')
     bechdel.head()
```

```
[7]:   year  title bechdel  budget  domgross  intgross  rated \
0  2013  21 & Over  FAIL  13.000000  25.682380  42.195766  other
1  2012   Dredd 3D  PASS  45.658735  13.611086  41.467257  other
```

2	2013	12 Years a Slave	FAIL	20.000000	53.107035	158.607035	R
3	2013	2 Guns	FAIL	61.000000	75.612460	132.493015	R
4	2013	42	FAIL	40.000000	95.020213	95.020213	PG-13

	imdb_rating	romcom	drama	action	sci-fi	runtime
0	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN	NaN
2	8.3	0.0	1.0	0.0	0.0	134.0
3	6.8	0.0	0.0	1.0	0.0	109.0
4	7.6	0.0	1.0	0.0	0.0	128.0

## 1(b)

movie title

## 1(c)

```
[8]: bechdel['title'].count()
```

```
[8]: 1794
```

1794 observations

## 1(d)

```
[8]: pg = bechdel['rated'].value_counts()['PG']
pg_13 = bechdel['rated'].value_counts()['PG-13']
r = bechdel['rated'].value_counts()['R']
other = bechdel['rated'].value_counts()['other']
print('parental guidance suggested -', pg, 'parents strongly cautioned -',
      pg_13, 'restricted -', r, 'other -', other)
```

parental guidance suggested - 257 parents strongly cautioned - 565 restricted - 691 other - 281

### 1.1.1 1(e)

In order for the sample to be representative of the population, it must be randomly selected and large enough. In this case, there is no information about whether the sample was randomly selected or not. In fact, it could have been quite otherwise, and many types of biases could have been involved, since the dataset uses only two websites as a source of information. Also, I wouldn't say it's big enough to represent the entire population of films. According to Google, there are about 500,000 films, which means that the sample is 0.3% (1794/500000) of the population. Therefore, I believe that the dataset is not representative of the movie set.

## 1(f)

```
[9]: bechdel = bechdel.dropna(subset = ['runtime'])
      bechdel.head()
```

```
[9]:   year          title bechdel  budget  domgross  intgross  rated \
2  2013      12 Years a Slave   FAIL    20.0  53.107035  158.607035    R
3  2013           2 Guns   FAIL    61.0  75.612460  132.493015    R
4  2013           42   FAIL    40.0  95.020213   95.020213  PG-13
5  2013       47 Ronin   FAIL   225.0  38.362475  145.803842  PG-13
6  2013  A Good Day to Die Hard   FAIL    92.0  67.349198  304.249198    R

      imdb_rating  romcom  drama  action  sci-fi  runtime
2             8.3     0.0    1.0     0.0     0.0    134.0
3             6.8     0.0    0.0     1.0     0.0    109.0
4             7.6     0.0    1.0     0.0     0.0    128.0
5             6.6     0.0    0.0     1.0     0.0    118.0
6             5.4     0.0    0.0     1.0     0.0     98.0
```

## 1.2 Question 2

### 2(a)

```
[10]: mod = smf.ols(formula = 'imdb_rating ~ budget + drama + sci-fi + romcom', data_
      ↪= bechdel).fit()
      mod.summary()
```

```
[10]: <class 'statsmodels.iolib.summary.Summary'>
      """

                                OLS Regression Results
=====
Dep. Variable:          imdb_rating    R-squared:                0.079
Model:                  OLS          Adj. R-squared:            0.077
Method:                 Least Squares    F-statistic:          34.04
Date:                  Thu, 08 Dec 2022    Prob (F-statistic):      2.77e-27
Time:                  10:39:16          Log-Likelihood:         -2131.6
No. Observations:      1591             AIC:                  4273.
Df Residuals:          1586             BIC:                  4300.
Df Model:               4
Covariance Type:       nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept         6.4645      0.046    140.073      0.000         6.374         6.555
budget             0.0012      0.000     2.635      0.009         0.000         0.002
drama             0.5445      0.049    11.198      0.000         0.449         0.640
sci-fi           -0.0378      0.071    -0.530      0.596        -0.178         0.102
romcom           -0.2111      0.086    -2.469      0.014        -0.379        -0.043
=====
Omnibus:                 82.261    Durbin-Watson:           1.857
Prob(Omnibus):            0.000    Jarque-Bera (JB):        109.843
Skew:                    -0.484    Prob(JB):                1.41e-24
```

Kurtosis: 3.849 Cond. No. 298.  
=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
"""

**2(b)**

Adjusted R<sup>2</sup>, 0.077.

**2(c)**

drama, romcom

**2(d)**

Holding all other variables constant, one unit increase in romcom is associated with -0.2111 units increase in the imdb ratings.

**2(e)**

100\*0.0012 = 0.12 units increase in the imdb ratings

### 1.3 Question 3

**3(a)**

We've seen that at 5% level, budget and sci-fi failed to be significant, and we failed to conclude that there is a relationship between those features and IMDB. Also, since our R-squared value is low, which implies that our data is not fit well into our model. Therefore, I would say that the four features would be a good predictor of the IMDB.

**3(b)**

```
[11]: x_df = bechdel[['budget', 'drama', 'romcom', 'sci-fi']]
      y_df = bechdel[['imdb_rating']]
      x_df.head()
```

```
[11]:    budget  drama  romcom  sci-fi
      2    20.0    1.0    0.0    0.0
      3    61.0    0.0    0.0    0.0
      4    40.0    1.0    0.0    0.0
      5   225.0    0.0    0.0    0.0
      6    92.0    0.0    0.0    0.0
```

```
[13]: x_df.head()
```

```
[13]:    budget  drama  romcom  sci-fi
      2    20.0    1.0    0.0    0.0
      3    61.0    0.0    0.0    0.0
```

4	40.0	1.0	0.0	0.0
5	225.0	0.0	0.0	0.0
6	92.0	0.0	0.0	0.0

### 3(c)

```
[12]: x_train, x_test, y_train, y_test = train_test_split(x_df,y_df,test_size = 0.2,
↳random_state=135)

x_train.head()
```

```
[12]:          budget  drama  romcom  sci-fi
570    40.043484    1.0    0.0    0.0
957    27.130243    0.0    0.0    0.0
692    28.088587    0.0    0.0    0.0
114   101.463857    1.0    0.0    0.0
1752   53.560590    1.0    0.0    0.0
```

```
[15]: len(x_train), len(x_test)
```

```
[15]: (1272, 319)
```

### 3(d)

```
[13]: regressor = LinearRegression()
regressor.fit(x_train, y_train)
print('Intercept', regressor.intercept_)           #we trained the algorithm! Now
↳print the intercept
print('Coefficient', regressor.coef_)
```

```
Intercept [6.42565479]
Coefficient [[ 0.00161455  0.58995507 -0.23433976 -0.0324223 ]]
```

### 3(e)

intercept: 6.4645 and 6.42565479 budget: 0.0012 and 0.00161455 drama: 0.5445 and 0.58995507  
sci-fi: -0.0378 and -0.23433976 romcom: -0.2111 and -0.0324223

the coefficient for drama in the sklearn regression model is higher than in q2.

### 3(f)


```
[14]: y_pred = regressor.predict(x_test)
for i in range(10):
    print(y_pred[i])
```

```
[6.42969275]
[6.47394669]
[7.02198122]
[6.48700856]
```

```
[6.4566541]
[7.04467172]
[6.45022752]
[6.53049602]
[6.46014901]
[6.45018468]
```

3(g)

```
[15]: compare = pd.DataFrame({'Actual IMDB rating': y_test.to_numpy().flatten(),
                             'Predicted IMDB rating': y_pred.flatten()})
plt.scatter(y_test,y_pred)
plt.xlabel('actual imdb')
plt.ylabel('predicted imdb')
plt.show()
```



3(h)

```
[19]: print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
                                     ↪y_pred)))
```

Root Mean Squared Error: 0.9284746856773564

RMSE tells me the average difference between the predicted and actual values. In this case, on average, the actual and predicted values differ by 0.9284746856773564.

3(i)

```
[20]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('R-squared', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 0.7351599189777014

Mean Squared Error: 0.8620652419436656

R-squared 0.024405470454573308

On average, IMDB shows that the IMDB score for a certain movie is 5, then the real IMDB score may be somewhere between ~4-6. I would not say that this is a very bad forecast, but also not the most accurate. In addition, we also have an R squared value of less than 3%, which makes me doubt the accuracy of the model. As a data scientist, I wouldn't use the model for important purposes.

## 1.4 Question 4

4(a)

```
[16]: bechdel['bechdel'].describe()
```

```
[16]: count      1591
      unique        2
      top        FAIL
      freq        892
      Name: bechdel, dtype: object
```

Fail: 892 observations

Pass: 1591-892=699 observations

4(b)

```
[22]: (892/1591)*100
```

```
[22]: 56.06536769327467
```

56.06536769327467% of movies fail the test

4(c)

```
[17]: x_df = bechdel[['year', 'budget', 'domgross', 'intgross',
                  'imdb_rating', 'romcom', 'drama', 'action', 'sci-fi']]
      x_df
```

```
[17]:   year  budget  domgross  intgross  imdb_rating  romcom  drama \
2   2013  20.000000  53.107035  158.607035         8.3     0.0    1.0
```



3	2013	61.000000	75.612460	132.493015	6.8	0.0	0.0
4	2013	40.000000	95.020213	95.020213	7.6	0.0	1.0
5	2013	225.000000	38.362475	145.803842	6.6	0.0	0.0
6	2013	92.000000	67.349198	304.249198	5.4	0.0	0.0
...	...	...	...	...	...	...	...
1788	1971	14.386286	70.780525	70.780525	6.2	0.0	0.0
1789	1971	305.063707	404.702718	616.827003	6.6	0.0	0.0
1790	1971	143.862856	59.412143	64.760273	7.6	0.0	0.0
1791	1971	12.659931	236.848653	236.848653	7.8	0.0	0.0
1793	1970	5.997631	53.978683	53.978683	6.2	0.0	0.0

	action	sci_fi
2	0.0	0.0
3	1.0	0.0
4	0.0	0.0
5	1.0	0.0
6	1.0	0.0
...	...	...
1788	1.0	1.0
1789	1.0	0.0
1790	0.0	0.0
1791	1.0	0.0
1793	0.0	0.0

[1591 rows x 9 columns]

#### 4(d)

```
[18]: scaler = StandardScaler()
# x_df[['year', 'budget', 'domgross', 'intgross']] =
scaler.fit_transform(x_df[['year', 'budget', 'domgross', 'intgross']])
x_df.head()
```

```
[18]:   year  budget  domgross  intgross  imdb_rating  romcom  drama  action  \
2  2013    20.0  53.107035  158.607035         8.3     0.0    1.0    0.0
3  2013    61.0  75.612460  132.493015         6.8     0.0    0.0    1.0
4  2013    40.0  95.020213   95.020213         7.6     0.0    1.0    0.0
5  2013   225.0  38.362475  145.803842         6.6     0.0    0.0    1.0
6  2013    92.0  67.349198  304.249198         5.4     0.0    0.0    1.0

      sci_fi
2      0.0
3      0.0
4      0.0
5      0.0
6      0.0
```

#### 4(e)

```
[19]: y_df = bechdel[['bechdel']]
      xtrain, xtest, ytrain, ytest = train_test_split(x_df,y_df,test_size = 0.2,
      ↪random_state=321)

      xtrain.head()
```

```
[19]:      year      budget      domgross      intgross      imdb_rating      romcom      drama  \
1203  2001  143.427209  415.208284  1167.441808           8.9         0.0         0.0
1599  1991   23.951765   59.415931   59.415931           6.6         0.0         1.0
585   2008   32.467689   68.368797  113.824188           7.3         1.0         1.0
1600  1990   71.319016  156.307443  434.511105           7.4         0.0         0.0
1551  1993  104.809829  135.525899  411.177020           6.3         0.0         0.0

      action      sci-fi
1203      1.0         0.0
1599      0.0         0.0
585      0.0         0.0
1600      0.0         1.0
1551      1.0         0.0
```

4(f)

```
[20]: xtrain
```

```
[20]:      year      budget      domgross      intgross      imdb_rating      romcom      drama  \
1203  2001  143.427209  415.208284  1167.441808           8.9         0.0         0.0
1599  1991   23.951765   59.415931   59.415931           6.6         0.0         1.0
585   2008   32.467689   68.368797  113.824188           7.3         1.0         1.0
1600  1990   71.319016  156.307443  434.511105           7.4         0.0         0.0
1551  1993  104.809829  135.525899  411.177020           6.3         0.0         0.0
...   ...         ...         ...         ...         ...         ...         ...
1612  1990   53.489262  215.222722  357.486568           7.6         0.0         0.0
909   2005   34.001444   7.517458   18.435915           7.1         0.0         1.0
140   2012  147.122592  104.926572  311.393491           7.3         0.0         0.0
616   2008  248.918952  183.300049  640.362486           6.7         0.0         0.0
1141  2002   55.692706   31.641581   57.808305           5.0         1.0         0.0

      action      sci-fi
1203      1.0         0.0
1599      0.0         0.0
585      0.0         0.0
1600      0.0         1.0
1551      1.0         0.0
...   ...         ...
1612      1.0         0.0
909      0.0         1.0
140      0.0         0.0
```

```
616      1.0      0.0
1141     0.0      0.0
```

```
[1272 rows x 9 columns]
```

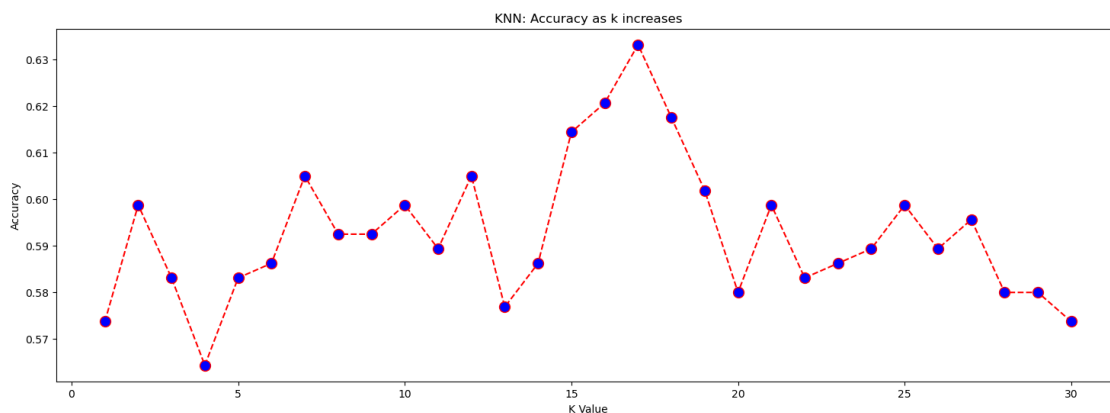
```
[21]: error = []
      accuracy = []

      for i in range(1,31):
          knn = KNeighborsClassifier(n_neighbors=i)
          knn.fit(xtrain, ytrain.to_numpy().flatten())
          pred_i = knn.predict(xtest)
          error.append(np.mean(pred_i != ytest.to_numpy()))
          accuracy.append(metrics.accuracy_score(ytest, pred_i))
```

```
[22]: # plot accuracy

plt.figure(figsize=(18, 6))
plt.plot(range(1, len(accuracy)+1), accuracy, color='red', linestyle='dashed',
        marker='o',
        markerfacecolor='blue', markersize=10)

# title and label axes
plt.title('KNN: Accuracy as k increases')
plt.xlabel('K Value')
plt.ylabel('Accuracy')
plt.show()
```



4(g)

From the graph,  $k = 17$  have the greatest accuracy.

```
[23]: classifier = KNeighborsClassifier(n_neighbors=17) #set
      ↪ k=5 (a common value for k)
      classifier.fit(xtrain, ytrain.to_numpy().flatten())
      ↪ #train algorithm! #was x_train, y_train
      ypred = classifier.predict(xtest)
      for i in range(6):
          print(ypred[i])
```

FAIL

PASS

PASS

PASS

FAIL

PASS

4(h)

```
[24]: from sklearn.metrics import classification_report, confusion_matrix
```

```
[26]: print(confusion_matrix(ytest, ypred))
```

```
[[131  50]
 [ 67  71]]
```

50 movies were incorrectly predicted, although they have passed the Bechdel test.

4(i)

```
[27]: print(classification_report(ytest, ypred))
```

	precision	recall	f1-score	support
FAIL	0.66	0.72	0.69	181
PASS	0.59	0.51	0.55	138
accuracy			0.63	319
macro avg	0.62	0.62	0.62	319
weighted avg	0.63	0.63	0.63	319

Recall for Pass is the portion of correct “positive” predictions for pass out of all positive predictions. It suggests that we predicted ‘Fail’ a little better than ‘Pass’.

4(j)

With an accuracy of 63%, I would say that our model worked pretty well. Thus, on average, 6 out of 10 films would have been predicted correctly. Although I would prefer it to be slightly increased, but I think the results are useful.